

Recurrent nets

Bryan Pardo


Deep Learning

Northwestern University


Deep Learning: Bryan Pardo, Northwestern University, Fall 2020

Dealing with time

- With a "standard" feed-forward architecture, you process data from within a window, ignoring everything outside the window.
- To get influence from the processing of earlier time steps, add nodes and connections
- This doesn't scale well




Let's look at that net.




- An entire new set of weights for EACH time step.
- Audio is sampled at 44,100 times per second
- The number of past time steps you could consider is limited by the architecture.
- The number of weights to learn quickly gets out of control.

Take an idea from CNNs and HMMs



- Markov property: The state of the world can be captured by knowing current state + immediately previous state
- Markov models use recurrent connections
- CNNs use the same set of shared weights on different parts of the input


Take an idea from CNNs and HMMs



- If all the windows share the same input weights (like in a feature map), then we only have the same number of weights as if we had a single window.
- This is a recurrent net.
- How do you train this?
- Are there any obvious limitations?


Backprop through time: “Unrolling”

- Pick a number of steps over which you’re going to “unroll” the net.
- Treat it like you’re training a convolutional neural net
- Pick the number of steps based on your frenemy: Exponential decay




Getting influence from the past: Skip connections (used in Highway networks)

- Widely used
- Limited by the length of the skip




Exponentially decaying influence

- If your network needs to connect information from a distant timestep, the influence of the earlier one tends to get lost
- Why? Exponential decay.



Exploding and vanishing gradients



- What if the weight onto the recurrent connection is greater than 1?
- What if the weight on the recurrent link is less than 1?
- What if it is exactly 1?

An RNN example: Language modeling

- In language modeling, the game is to be able to predict the next word, given the previous N words.
- Examples
 - “Two plus two equals...”
 - “A stitch in time saves....”
 - “I never did...”


Our text encoding

- 1000 most common English words. + start + stop + other
- Encoding: 1003 element one-hot vector for each word in a sentence
 - Word index determined by popularity
 - Start = 1001
 - Stop = 1002
 - Other (any word not in the top 1000) = 1003
- Examples:
 - An apple is good for you. -> [1001, 48, 927, 121, 7, 26, 1002]
 - Lilliputian dilatants prognosticate parsimoniously! -> [1001, 1003, 1003, 1003, 1003, 1002]

The goal: predict the next token

- Each sentence is its own label.
- Given “An apple is...”, predict “good” as the next word.
- Our model output will be a probability distribution over the 1003 element vector (top 1000 words + start + start + other).
- We can use cross-entropy loss, comparing the one-hot vector to the probability vector output by the model.

Our network



A RNN with 4 hidden nodes : how many weights?

OUTPUT a 1003 element probability distribution over the set of words.

Output layer $[0.01, .0000098, \dots, .0023, \dots, .001, 0, .000053]$ Softmax activation



ALL HIDDEN NODES ARE FULLY CONNECTED TO THE OUTPUT LAYER

Hidden layer



Sigmoid activation



ALL HIDDEN NODES ARE FULLY CONNECTED TO THE INPUT LAYER

Input +
prev state

$[0, 0, 0, \dots, 1, \dots, 0, 0, 0]$

INPUT word $w(t)$: a 1003 element
one-hot vector encoding word t .

$[.2, 0, .001, .3]$


Previous state $s(t-1)$: a vector of the output
from each hidden unit from time $t-1$

What if we use a more realistically sized net?

- Dictionary size = 50,000
- Hidden states = 100
- $50,000 * 100 * 2 = 10,000,000$
- It's just that easy to have 10 million weights.
- Adding a couple of extra hidden layers (even fully connected ones) doesn't cost you much, compared to the dictionary size.

RNN: the math

OUTPUT a 1003 element probability distribution over the set of words.



RNN: the math

OUTPUT a 1003 element probability distribution over the set of words.

Output layer

$$y_k(t) = g \left(\sum_k s_j(t) v_{jk} \right)$$

$$g(z_m) = \frac{e^{z_m}}{\sum_k e^{z_k}}$$

Hidden layer

$$s_j(t) = \sigma \left(\sum_i u_{ij} x_i(t) \right)$$

$$\sigma(z) = (1 + e^{-z})^{-1}$$

Input +
prev state

$$w(t)$$

$$s(t-1) = [s_1(t-1), \dots, s_n(t-1)]$$

INPUT word $w(t)$: a 1003 element
one-hot vector encoding word t .

RNN: making a prediction

$$prediction = \hat{w}(t + 1) = \operatorname{argmax}_k [y_1(t), \dots, y_k(t), \dots, y_m(t)]$$

Output layer

$$y_k(t) = g\left(\sum_k s_j(t)v_{jk}\right)$$

$$g(z_m) = \frac{e^{z_m}}{\sum_k e^{z_k}}$$



Hidden layer

$$s_j(t) = \sigma\left(\sum_i u_{ij} x_i(t)\right)$$

$$\sigma(z) = (1 + e^{-z})^{-1}$$



Input +
prev state

$$x(t) = w(t) + s(t - 1)$$

$$w(t)$$

$$s(t - 1) = [s_1(t - 1), \dots, s_n(t - 1)]$$

INPUT word $w(t)$: a 1003 element
one-hot vector encoding word t .

Perplexity

- A measure of how hard it is to guess the next word.
- The exponentiation of the cross-entropy

$$\text{Perplexity} = 2^{H(p)} = 2^{-\sum_x p(x) \log_2(q(x))}$$

- A commonly used measure of how well a language model is doing
- Measures how confused the model is (how many choices it has reduced the next word to)

A language model is a generative model

- If you have something that predicts the next word, you have something that can “generate” the next word.
- Sentence completion is possible
- Sentence generation is possible

Language model as a generative model:


generation = $\hat{w}(t + 1)$ = sample from: $[y_1(t), \dots, y_k(t), \dots, y_m(t)]$

Output layer 	$y_k(t) = g\left(\sum_k s_j(t)v_{jk}\right)$	$g(z_m) = \frac{e^{z_m}}{\sum_k e^{z_k}}$
Hidden layer 	$s_j(t) = \sigma\left(\sum_i u_{ij} x_i(t)\right)$	$\sigma(z) = (1 + e^{-z})^{-1}$
Input + prev state 	$x(t) = w(t) + s(t - 1)$	
	$w(t)$	$s(t - 1) = [s_1(t - 1), \dots, s_n(t - 1)]$


INPUT word $w(t)$: a 1003 element one-hot vector encoding word t .

Recurrent neural network-based language model (Interspeech 2010)


Generating a new sentence with the model




Generating a new sentence with the model



Generating a new sentence with the model



Generating a new sentence with the model



Getting more context

- We predict/generate a new token, based on a prior sequence.
- Our generated output is contextually informed by the past
- But wait....if our training data is whole sentences, can't we do the same thing from the "future" (i.e. the next word or rest of sentence)?
- Sure we can. Just feed in the sequence backwards.

RNN: predicting the “past” based on the “future”

$$prediction = \widehat{w}(t - 1) = \max_k [y_1(t), \dots, y_k(t), \dots, y_m(t)]$$

Output layer



$$y_k(t) = g\left(\sum_k s_j(t)v_{jk}\right)$$

$$g(z_m) = \frac{e^{z_m}}{\sum_k e^{z_k}}$$

Hidden layer

$$s_j(t) = \sigma\left(\sum_i u_{ij} x_i(t)\right)$$

$$\sigma(z) = (1 + e^{-z})^{-1}$$

Input +
next state

$$x(t) = w(t) + s(t + 1)$$


$$w(t)$$

$$s(t + 1) = [s_1(t + 1), \dots, s_n(t + 1)]$$

INPUT word $w(t)$: a 1003 element
one-hot vector encoding word t .

Bidirectional RNN

- Inform output layer's probability distribution using a forward layer and a backwards layer
- The generated token(s) are influenced by both previous and subsequent context



Multi-layer RNN

- You can have multiple hidden layers, where layer n feeds into layer $n+1$




Image from Graves, Alex. "Generating sequences with recurrent neural networks." *arXiv preprint arXiv:1308.0850* (2013).

Long-Short Term Memories

Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." *Neural computation* 9.8 (1997): 1735-1780.

Here's a problem. What can learn to do it?

- X is a finite-length sequence composed of tokens, where each token $x_n \in \mathbb{R} \cup \{a, b\}$.
- The length of X is unknown.
- Before beginning, the total = 0.
- Iterate through X and do the following
 - If $x_n = a$, add x_{n+1} to the total.
 - If $x_n = b$, return the total and reset the total to 0.

Let's play


$$a, -1, a, 100, b = 99$$

$$1, 3, -5, a, 5, -1, 8, 2, 0, a, 9, b = 14$$

$$1, 3, a, -5, 0, 0, 0, 0, a, 5, -1, 7, 2, 0, a, 9, -12, b = 9$$


Feed-forward: Fixed-length time window

- If your network needs to connect information from outside the window, you lose.



RNN: exponentially decaying influence

- If your network needs to connect information from a distant timestep, the influence of the earlier one tends to get lost
- Why? Exponential decay.



Long Short Term Memory Units (LSTMs)

- Added a way of storing data over many time steps without decay
- Let networks to handle problems with long term dependencies




Image from: Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." *Neural computation* 9.8 (1997): 1735-1780.

LSTM training

- Error is propagated indefinitely through its memory cell, the constant error carousel (CEC)
- Error flow back through the unit is truncated at the incoming weights.





Image from: Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." *Neural computation* 9.8 (1997): 1735-1780.


An easy-to-follow-visual of a modern LSTM

Forget gates introduced in: Gers, F.A.
et al (1999). "Learning to forget:
Continual prediction with LSTM". 9th
*International Conference on Artificial
Neural Networks*: 1999




Amazing gifs from Michael Phi's <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

Forget Gate




Amazing gifs from Michael Phi's <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

Input Gate




Amazing gifs from Michael Phi's <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

Cell State




Amazing gifs from Michael Phi's <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

Output Gate



Amazing gifs from Michael Phi's <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

The math of the modern LSTM



$$i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1} + W_{ci}c_{t-1} + b_i)$$
$$f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1} + W_{cf}c_{t-1} + b_f)$$
$$c_t = f_t c_{t-1} + i_t \tanh(W_{xc}x_t + W_{hc}h_{t-1} + b_c)$$
$$o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1} + W_{co}c_t + b_o)$$
$$h_t = o_t \tanh(c_t)$$

Image adapted from Graves, Alex. "Generating sequences with recurrent neural networks." *arXiv preprint arXiv:1308.0850* (2013).

Forget gates introduced in: Gers, F.A. et al (1999). "Learning to forget: Continual prediction with LSTM". *9th International Conference on Artificial Neural Networks*: 1999

How many weights for a single LSTM unit?

Input gate $i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1} + W_{ci}c_{t-1} + b_i)$

forget gate $f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1} + W_{cf}c_{t-1} + b_f)$

memory $c_t = f_t c_{t-1} + i_t \tanh(W_{xc}x_t + W_{hc}h_{t-1} + b_c)$

output gate $o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1} + W_{co}c_t + b_o)$


final output $h_t = o_t \tanh(c_t)$

$$\text{number of weights} = 4|x| + 4|h| + 3 + 4$$

How many weights for this network?

- Input: 50,000 word vocabulary, 4 LSTM layers of 100 cells per layer
- Compare that to a vanilla RNN with the same number of layers and vocabulary size....
- Can we shrink closer to a vanilla RNN but keep advantages of an LSTM?

GRU: A simplified LSTM



Images from Michael Phi's <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

GRU: The Math

A linear interpolation between previous output and candidate output

Final output
$$h_t = (1 - z_t)h_{t-1} + z_t \tilde{h}_t$$

Determines how much to make the output be influenced by the previous hidden state vs the current input.

Update gate
$$z_t = \sigma(W_z \mathbf{x}_t + U_z \mathbf{h}_{t-1})$$

Vector of all reset gates in the hidden layer

Candidate output
$$\tilde{h}_t = \tanh(W_{\tilde{h}} \mathbf{x}_t + U_{\tilde{h}} (\mathbf{r}_t \odot \mathbf{h}_{t-1}))$$



Vector of all outputs in the hidden layer

Reset gate
$$r_t = \sigma(W_r \mathbf{x}_t + U_r h_{t-1})$$

Math based on: Chung, Junyoung, et al. "Empirical evaluation of gated recurrent neural networks on sequence modeling." *arXiv preprint arXiv:1412.3555* (2014)

LSTM/GRU Plusses and Minuses

- Lets networks handle problems with long term dependencies
- This lets LSTMs (or GRU) solve problems simple recurrent architectures cannot
- Still has trouble with XOR (time-delayed XOR where you XOR two inputs that are an unknown number of time steps apart)
- Lots of extra weights compared to regular cells
- Long and slow to train
- Not easy to inspect networks to understand them